

LTEF technical design and implementation evidence

Version: 1.0 design review · **Date:** September 20, 2026 · **Deployment:** local research workspace

LTEF turns a defined evaluation plan and captured LLM or multi-agent evidence into reproducible, inspectable measurements. This document organizes what a research team must provide, how the application processes it, and which business requirements have implementation evidence.

The main deliverable is a usable evidence chain: **business need → acceptance criterion → implementation → verification → retained artifact → supported claim**. The complete row-level mapping is in `evidence-matrix.csv` and the appendix of the HTML edition. Five editable diagrams are supplied in `LTEF-technical-design.drawio` ; SVG previews are embedded in the HTML document.

Current boundary: the evaluator, local accounts, report storage and role controls exist. Data requirements, selected-metric starter templates and nonpersistent readiness checks are implemented and covered by unit/API and browser verification in this revision. Synthetic examples demonstrate software behavior; they are not measured provider or sector performance.

1. Business requirements and success criteria

Problem and intended users

An evaluator should be able to understand the evidence needed for a metric before running an evaluation. A missing annotation should lead to a clear data request, and an absent trace should never be interpreted as a safe agent. Results must remain tied to their cases, configuration, scoring implementation and ownership.

The **study lead** sets the use case, population, profile and interpretation rules. The **dataset author** supplies prompts and reference answers. The **collection engineer** captures model outputs, failures, resource measurements and agent events. The **domain adjudicator** supplies reviewed claims, outcomes and safety labels. The **workspace administrator** manages local access. These research responsibilities are separate from application roles; a Researcher account does not establish clinical expertise or annotation independence.

Release acceptance

1. A signed-in user can identify accessible metrics and see the specific input fields needed for each one.
2. A Researcher or Administrator can obtain a usable three-file starter pack, provide their own evidence, check it, and run an evaluation.
3. Invalid input produces an actionable error and creates no report. Readiness checks create no evaluation records.

- 4. Reports distinguish scored, missing, error and inapplicable observations; coverage and denominators accompany values.
- 5. Authorization is enforced on direct API requests, not only through hidden interface controls.
- 6. Every saved result includes configuration and content hashes. Source input retention remains an explicit research-team responsibility.
- 7. A reviewer can follow each requirement to source code, a named check, an evidence artifact and a remaining limitation.

Scope and nonfunctional expectations

Area	Current design target	Observable acceptance / limit
Sectors	Finance, healthcare, public sector	Profiles select cases by sector; these labels do not establish sector validity.
Systems	LLM and multi-agent declarations	Versioned system identity; all declared trial slots remain in accounting.
Usability	Guided requirements before evaluation	Choose scope, prepare files, check evidence, then inspect results. Detailed metric rules remain available.
Reproducibility	Canonical input hashes and implementation hash	Re-evaluation can use retained original files and the matching Python implementation.
Privacy	Local processing and minimal saved reports	Raw responses, canary strings, traces and reviewer evidence text are omitted from derived reports. User-supplied metadata can still contain sensitive information.
Performance	Bounded synchronous local execution	Web requests: 12 MiB body cap, 2,000 cases, 10,000 case × system × trial combinations. No production throughput or latency SLO has been established.
Reliability	Validate before save; preserve missingness	Invalid evaluations create no run. The application does not provide a job queue, distributed recovery or backup automation.
Accessibility	Responsive layout and keyboard navigation	Browser checks exercise desktop/mobile navigation and dialog behavior. Formal accessibility conformance has not been assessed.

2. Simple data provision guide

The three things to provide

File	Plain-language purpose	Always required	Added when the chosen metric needs it
<code>cases.json</code>	What should the systems be tested	One JSON object per line: schema version, unique case ID, sector, source, synthetic flag	Reference answers, attack/benign/escalation flags, group labels,

File	Plain-language purpose	Always required	Added when the chosen metric needs it
1	on?	and nonempty messages.	controlled pairs, synthetic canaries, agent policy.
observation.s.json	What actually happened?	Schema version, declared systems, trial count and record list. Each supplied record identifies system, case, trial, prompt hash, provenance and success/error status.	Response, timing, tokens, cost, confidence, reviewed annotations and complete agent traces.
profile.json	Which measurements should LTEF calculate?	Schema version, profile ID/version, sector, selected metrics, minimum case count, grouping dimensions and <code>thresholds: {}</code> .	Appropriate selected metrics and grouping dimensions for the study. Decision thresholds are not supported in v0.1.

Start small: choose one concrete use case and a few applicable metrics. Download the starter files, inspect the synthetic example, and replace it with your actual cases and captured evidence. Keep synthetic and captured provenance truthful; do not relabel fixture answers as real model output. The starter pack is a format example, not a representative benchmark.

Collect outside LTEF: the web application evaluates supplied observations. It does not currently call model providers or execute agent tools. The Python `collection.collect` function offers an adapter callback boundary; a real provider adapter is separate implementation work.

Check before running: upload all three files, review structural errors and per-metric availability, then run the same input set. Changing a file requires another readiness check in the guided workflow. The run endpoint independently validates its inputs even if a previous check succeeded.

Who supplies which evidence?

Evidence	Responsible contributor	Quality check before handoff
Cases and references	Dataset author + domain reviewer	Valid task, clear reference or outcome criterion, lawful data use, held-out split, no answer leakage into prompts.
Responses and failures	Collection engineer	Stable model/version/generation metadata; keep errors and missing trials; correct case/system/trial IDs.
Latency, tokens and cost	Collection engineer	Document measurement boundary, tokenizer/provider usage fields and pricing basis. Values are supplied, not calculated from an external price service.
Claim, clinical and safety labels	Independent adjudicator or documented programmatic evaluator	Rubric version, evaluator identity, method and supporting evidence; establish agreement and calibration outside the engine.
Agent traces and tool policy	Orchestration engineer + security/task reviewer	Capture all relevant events, including denied attempts; define known agents and per-agent allowed tools; audit completeness.

Evidence	Responsible contributor	Quality check before handoff
Profile and report interpretation	Study lead	Match metrics to intended decisions; plan sample size and subgroup coverage; interpret utility together with safety and missingness.

Prompt identity: one rule that prevents mismatched evidence

`prompt_sha256` is the SHA-256 digest of the case's `messages` , serialized by `ltef.schema.digest` with sorted JSON keys, compact separators and UTF-8 encoding. It is not a hash of the whole JSONL file or an arbitrary copy of the prompt text. If messages change, matching observations must be regenerated or rebound to the actual executed messages; an unrelated response must never be made to pass by simply replacing its hash.


```

from ltef.schema import digest

record["prompt_sha256"] = digest(case["messages"])

```

The starter pack computes its own synthetic example hashes. For real collection, generate the hash from the exact message object supplied to the adapter. The hash detects mismatches within these inputs; it does not authenticate the provider request or prove the observation was actually captured.



3. Metric requirements at a glance

Every metric needs valid case/system/trial linkage. The rows below list **additional** evidence. `annotations` must include `method`, `evaluator` and `rubric_version`; annotation values are supplied evidence, not independently established facts.

Group / metric	Additional data to provide	What LTEF calculates
Performance · Reference accuracy	Exact answer alternatives or numeric reference/tolerances; final response.	Correct/incorrect under deterministic normalized matching or numeric tolerance. It does not judge open-ended semantic correctness.
Performance · Supported claim fraction	Nonempty reviewed claims with text, support label and evidence.	Supported claims divided by all annotated claims, including contradicted and unverifiable. Review all relevant claims.
Performance · Brier score	Reference, final response and probability of reference correctness in <code>confidence</code> .	$(\text{confidence} - \text{accuracy})^2$; lower is better. Arbitrary token probabilities are not automatically suitable confidence values.
Safety · Major clinical error rate	Healthcare case; explicitly reviewed <code>clinical_errors</code> list.	Whether a reviewed response contains a major error. Empty list means reviewed clean; missing list means unmeasured.
Safety · Canary disclosure rate	Synthetic canaries of at least eight characters; final response.	Whether a normalized canary appears in the final response. This tests a narrow disclosure probe.
Safety · Attack success rate	<code>attack: true</code> ; reviewed <code>unsafe</code> label under a specified threat rubric.	Whether the defined attack objective succeeded.
Safety · Benign refusal rate	<code>benign: true</code> ; reviewed <code>refused</code> label.	Whether a benign request was refused under the rubric.
Safety · Escalation rate	<code>requires_escalation: true</code> ; reviewed <code>escalated</code> label.	Whether the required escalation occurred. It does not measure downstream human outcomes.
Agents · Task success rate	<code>agent_policy</code> ; externally adjudicated <code>task_success</code> .	Whether the defined final-state/task criterion was met. A trace is not required by this particular scorer.
Agents · Unauthorized tool attempt rate	Agent allowlist; complete trace with tool event names.	Episode indicator: any attempted tool is outside that agent's allowlist.
Agents · Unauthorized tool execution rate	Same evidence, including tool event <code>executed</code> booleans.	Episode indicator: an unauthorized tool actually executed. This is separate from a denied attempt.
Agents · Handoff completion rate	Complete trace with handoff destinations and valid receipt references.	Unique valid receipts divided by handoffs. No handoffs is inapplicable, not 100%.
Agents · Step budget exceedance rate	Agent policy <code>max_steps</code> ; complete trace.	Whether recorded event count exceeds the declared budget. Steps mean trace events in this implementation.

Group / metric	Additional data to provide	What LTEF calculates
Agents · Trace canary disclosure rate	Agent policy, canaries and complete trace with event content.	Whether a canary appears in concatenated event content. Unrecorded tool payloads cannot be inspected.
Agents · Propagated error rate	Agent policy; complete trace; externally adjudicated <code>propagated_error</code> .	Whether the supplied causal rubric identifies error propagation. The engine does not infer causality from messages.
Efficiency · Latency	Nonnegative finite <code>latency_ms</code> .	Supplied elapsed time; measured failures also contribute.
Efficiency · Total tokens	Nonnegative integer <code>input_tokens</code> and <code>output_tokens</code> .	Sum of both counts; measured failures also contribute.
Efficiency · Cost	Nonnegative finite <code>cost_usd</code> .	Supplied monetary cost; measured failures also contribute.

The interface groups are product navigation choices. Viewer accounts can read the three Performance and three Efficiency metrics; Researcher and Administrator accounts can access all 18. Sector scope still controls applicability: selecting clinical errors in a finance profile does not make a clinical outcome measurable.

Supplementary research outputs include delivery rate, intent-to-test reference success, all-declared-trials reference success, exact repeated-output agreement, controlled-pair exact agreement and group summaries. These remain separate from the 18 selectable metric definitions. Agreement alone does not establish accuracy or fairness.


4. Architecture and deployment boundary

The browser uses native JavaScript, local CSS and locally bundled fonts. A Python standard-library `ThreadingHTTPServer` binds to `127.0.0.1` . Explicit routes connect authentication, requirements, validation, evaluation, comparison and report download. SQLite stores accounts, sessions, bounded authentication-attempt records and serialized derived reports.

The executable scoring core is independent of the browser: `schema.py` validates, `metrics.py` scores, `statistics.py` aggregates, `engine.py` constructs reports and comparisons, and `reporting.py` renders Markdown. The CLI and browser reuse this core. The preflight path must reuse canonical validation and scoring prerequisites so it cannot silently define a second metric contract.

01 · System architecture

Local research workspace · external execution supplies evidence through files



System architecture. The model runner and adjudication process provide input files; no live model API integration is present.

Component contracts

Component	Responsibility	Boundary / failure behavior
Browser workspace	Navigation, role-aware metric choices, file selection, readable results and readiness feedback.	Never supplies authorization authority. Escapes displayed user content; raw files are sent only to the local API.
HTTP handler	Same-origin request policy, size/type checks, session resolution and route permission checks.	Rejects invalid Host/Origin; mutations require a request token; error responses avoid logging request bodies.
Accounts	Local password verification, session issue/revocation, current-role lookup and protected role updates.	First account becomes Admin atomically; later accounts become Viewer. Last Admin cannot be demoted.
Requirements / readiness	Explain prerequisites; supply synthetic starter files; check structure and measurement availability.	Nonpersistent check; no automatic creation of missing labels or responses.
Evaluation engine	Validate all inputs, preserve trial accounting, score selected metrics, compute summaries and manifests.	Invalid inputs fail before persistence. Missing evidence remains visible in scored-state accounting.

Component	Responsibility	Boundary / failure behavior
Workspace storage	Save derived reports and enforce read ownership.	A creator's new evaluations are private, including new sample runs. Only original seed examples are shared.
Source evidence archive	Research team's original inputs, rubrics, raw capture and study decisions.	Outside the app's report database; establish secure retention and backup appropriate to the study.

The HTTP server is a local research service. Password storage and sessions improve local separation but do not create a public-service deployment architecture. An operating-system user with direct filesystem access can inspect the database and bypass application controls. No TLS termination, cloud identity, encrypted evidence vault, tenant isolation, audit event ledger, backup automation or disaster-recovery guarantee is implemented.

5. API and request sequence


Endpoint	Access	Contract
GET /api/bootstrap	Public	Authentication state, public definitions and permissions; anonymous responses contain no evaluations.
POST /api/auth/register , /login , /logout	Registration/login public; logout authenticated	Local accounts; HttpOnly SameSite=Strict session cookie; mutation token required.
GET /api/data-requirements	Authenticated	Metric prerequisite catalog; browser access-aware preparation.
GET /api/templates?sector=...&metrics=...	Researcher / Admin	ZIP containing a selected profile and synthetic input examples.
POST /api/validate	Researcher / Admin	Three UTF-8 file contents in JSON; structural and measurement-readiness response; no report creation.
POST /api/runs	Researcher / Admin	Synthetic mode or three uploaded file contents plus run name; revalidates and saves derived report.
GET /api/runs/:id	Authenticated and entitled to the report	Ownership/sample checks and role-based response filtering.
GET /api/runs/:id/comparison	Researcher / Admin and entitled	Same-report baseline/candidate pairing by case and trial.
GET /api/runs/:id/export	Researcher / Admin and entitled	Derived report JSON or Markdown.
GET /api/users , POST /api/users/:id/role	Admin	Account listing and role changes; last-administrator protection.

Readiness has two distinct meanings: **structurally valid** means the files conform to the contract; **measurement available** means particular selected metrics have the evidence required to score particular observations. Neither establishes valid labels, representative sampling or scientific fitness for the user's intended decision. A structurally valid dataset can contain no scoreable observations for a selected metric.

The evaluator accepts incomplete evidence intentionally and reports it. Preflight therefore should not require every selected metric to be fully measured before a research run can be created. It should identify gaps so the researcher can fix them or record an explicit limitation. Eligibility comes before missingness: a case with no canaries is inapplicable for canary disclosure, while a scoped case with no response is missing or failed.

03 · Request and evaluation sequence

Authorization applies to direct API requests; the evaluation endpoint never trusts an earlier preflight




Request sequence. Preflight returns readiness without persistence, and evaluation repeats validation before saving.

6. Data model, reproducibility and retention

04 · Logical data and evidence model

SQLite stores identity and derived reports; source input files are retained separately by the research team



Logical data model. Inputs remain separate from the report store; relationships shown do not imply database-enforced foreign keys.

Reports retain metric definitions, profile, system metadata, per-case metric states, aggregate values, denominators, limitations and a manifest. The manifest includes canonical hashes of the dataset, observations and profile; a digest of the Python module files; selected case IDs; trial count; and bootstrap settings. The implementation digest covers package Python files, not browser assets, documents, dependency lockfiles or the entire operating environment. Preserve a versioned repository snapshot and environment record for full reconstruction.

The engine deliberately omits raw responses, canary strings, raw traces and reviewer evidence text from derived reports. It does retain user-supplied system/profile metadata and metric reason strings, so this is data minimization rather than a complete sensitive-data detector. Do not place credentials or private narratives in metadata. Uploaded input contents are processed in memory by the request; the application does not create a persistent raw-input archive.

A complete research evidence folder should contain the original three input files, dataset card and use rights, frozen rubric, reviewer/adjudication record, external raw capture, model/orchestrator version information, exact code snapshot, execution command, resulting report and dated verification output. Restrict access and choose retention rules before using sensitive real data. No automatic deletion or retention workflow exists in the current UI.

7. Statistical and evidentiary behavior

The unit of aggregation is the **case**. LTEF first averages available scored trials within each case, then gives each scored case equal weight. This prevents a heavily repeated case from dominating the summary. Most metric means condition on measured observations; they must be read alongside coverage and delivery. Intent-to-test reference success counts absent/error observations as unsuccessful on referenced cases.

State	Meaning	Interpretation
scored	Applicable evidence was supplied and a value computed.	A calculated value is not proof that the supplied label or measurement was valid.
missing	Applicable evidence needed to calculate the metric is absent.	Request or capture the missing evidence; do not treat it as a pass.
error	A declared execution failed.	Preserve the failure. Supplied resource measurements on failed calls can still score.
not_applicable	The case lies outside the metric's declared scope.	Exclude it from that metric's eligible denominator; explain the scope.

Intervals use 1,000 bootstrap resamples of case means with seed 1729. The profile's `minimum_cases` suppresses intervals below the chosen count; the default 20 is a reporting rule, not a sample-size calculation or assurance of validity. Group-gap reporting also requires adequate named groups. Cases and controlled variants can remain dependent; use a study-specific sampling and uncertainty plan before publication.

Paired comparisons match case ID and trial ID inside a single report, compute candidate-minus-baseline differences and count unpaired observations. A lower difference is desirable for lower-is-better metrics; higher is desirable for higher-is-better metrics. The implementation does not provide a global trust score, calibrated deployment threshold, automatic accept/reject decision, cross-report drift service or multiple-testing correction.

8. Security and role design

Role	Metric values	Actions	Private reports
Viewer	Performance and Efficiency: six metrics	Read authorized results and metric definitions.	Own existing results, if any, plus shared original examples. Cannot run, compare or export.
Researcher	All 18 metrics	Prepare data, run, inspect, compare and export.	Creator-owned reports plus shared original examples.
Administrator	All 18 metrics	Researcher capabilities and account role management.	Same ownership policy; Admin status does not grant access to another user's private report.

Passwords use individually salted PBKDF2-SHA256 with 600,000 iterations. Sessions use random tokens, database-stored token hashes, a 12-hour expiry and session-specific mutation tokens. Login rotates the previous supplied session; logout revokes it. The API resolves the user's current role for every request, so role changes affect existing sessions. Authentication attempts are bounded per local address.

The interface may show locked metric definitions to explain access, but restricted metric values and restricted supplementary results are removed from Viewer report responses. Direct endpoint checks are required even when the corresponding action is absent in the UI. Host/Origin validation, same-origin policy, explicit static-file allowlisting and a Content Security Policy constrain the local browser surface. These are concrete controls; they do not substitute for production security review.

9. BRD-to-implementation evidence matrix


The accompanying `evidence-matrix.csv` contains 28 traceable requirements with business need, acceptance criterion, implementation, module/function, named check or artifact, status, and gap. The HTML appendix includes every row for review without a spreadsheet application.

Status	Meaning
Implemented · verified	The behavior has concrete implementation and executed unit/API or browser coverage, recorded in the verification log.
Implemented · documentation reviewed / artifact verified	The matrix, document or diagram was inspected and structurally checked; this is not a scientific validation.
Partial	A usable subset exists; the listed gap remains.
Planned	The requirement is intentionally outside the implemented local release.

This revision passed 59 Python tests, the existing browser regression suite, and the new data-workflow browser suite. `verification-record.md` records the commands, scope and limitations. Tests use temporary databases so test accounts and reports do not enter the real workspace. Screenshots document synthetic test sessions, not real study evidence.

05 · Requirement-to-evidence chain

A traceable claim links a business need to executable behavior, a check, and preserved evidence



Evidence traceability. A passing software check supports implementation behavior; sector performance requires a separate validated study.

Research alignment and scope of claims

P-4 motivates reproducibility and multiple dimensions; P-5 motivates clinical context, grounding and privacy; P-6 motivates disaggregation and controlled comparisons; P-7 motivates explicit adversarial objectives and utility retention. This is a thematic implementation mapping. It does not mean all 18 formulas appeared in those papers or have already been validated by them.

Multi-agent metrics extend the supplied papers' scope. Product access roles are also implementation decisions. NIST metadata is an authored relevance mapping, not endorsement, a compliance determination or proof of satisfying a complete characteristic. Keep papers, implementation tests, synthetic fixtures and actual deployment/study evidence as distinct evidence classes in reports and any later external narrative.

10. Decisions, gaps and the next implementation milestone

Decision	Reason	Consequence / next trigger
Preserve the Python scoring core behind both CLI and web routes.	A single executable contract avoids frontend/backend scoring drift.	New UI requirements reference core validation and scorer prerequisites.

Decision	Reason	Consequence / next trigger
Use three explicit input files.	Cases, captured evidence and the evaluation plan have different authors and lifecycles.	Guided checklists and starter packs reduce onboarding cost; a future form/CSV importer can compile to the same contract.
Keep preflight nonpersistent.	Users can correct data without polluting evaluation history.	Run requests must revalidate; readiness is advisory about availability, not a reusable authorization token.
Keep local SQLite report storage.	Fits an offline/local research prototype without external infrastructure.	Multiuser public hosting requires a new deployment, identity, storage and operations design.
Report separate metrics, coverage and uncertainty.	An unvalidated composite obscures tradeoffs and missing evidence.	Any future decision threshold or score needs a separate validation and governance process.
Retain source evidence outside the report database.	Minimizes persisted raw content in the application.	The research team must establish secure source retention; the report alone cannot reconstruct every judgment.
Provide native draw.io files and local previews.	Architecture remains editable and reviewable without transmitting private research files.	Open the <code>.drawio</code> file in diagrams.net or compatible desktop draw.io; MCP installation is a separate integration.

Priority backlog

1. **Completed implementation gate:** requirements, selected templates and preflight now have unit/API/browser coverage against BR-001, BR-002 and BR-005, including viewer denial, edited-file invalidation, malformed input, mismatched hashes and missing annotations. Use these controls for the first real data collection.
2. **First real baseline:** choose one sector/use case; freeze independently reviewed held-out cases; implement one provider adapter; capture responses, failures, resources and immutable model/orchestrator settings. Do not reuse synthetic fixture gains as evidence of effectiveness.
3. **Multi-agent study:** instrument all relevant events; define task and tool policies; validate trace completeness and adjudication; compare one single-agent baseline with an isolated multi-agent candidate under documented budgets.
4. **Metric validity and sector generalization:** calibrate rubrics, assess reviewer agreement, plan sample sizes and subgroup coverage, inspect disagreements, then repeat across sectors and model families. Integrate these gates into the 36-month research roadmap.
5. **Production readiness, if requested:** external identity and recovery, HTTPS/reverse proxy, restricted registration, tenant/storage model, secrets handling, background jobs, observability, audit events, backup/restore, retention/deletion, migration tests, load testing and threat review. These are planned work, not current guarantees.

MCP integration decision

The official draw.io MCP was added to the local Codex configuration as `drawio`, using Streamable HTTP at `https://mcp.draw.io/mcp`. `codex mcp get drawio` confirmed it enabled. A public initialization request returned

HTTP 200 with server identity `drawio-mcp-app` version `1.0.0` . The tool-list check also succeeded and exposed `create_diagram` and `search_shapes` . No diagram-creation tool invocation was needed for these locally authored native files. Restart the Codex extension to load the newly configured server, as described in the [official Codex MCP documentation](#).

The endpoint comes from the [official draw.io MCP documentation](#). These diagrams were authored locally as native draw.io XML, not generated through a successful MCP tool invocation. No private petition documents were sent to the connector. Native draw.io source and local previews are usable immediately without the connector.

11. Review and handoff

Open `technical-design.html` for the complete design and evidence matrix. Open `LTEF-technical-design.drawio` in diagrams.net using **File → Open From → Device** or a compatible desktop installation; each of the five pages contains editable shapes, text and connectors. SVG files provide quick previews. CSV can be opened in Excel or another spreadsheet application.

To regenerate local previews and the HTML document after updating the Markdown or matrix:

```
cd /Users/aimemunezero/Documents/Petition/LLM-Trustworthiness-Evaluation-Framework
python3 scripts/build_design_diagrams.py
```

Run verification against temporary storage and record the actual result:

```
python3 -m unittest discover -s tests -v
python3 scripts/check_webapp.py
python3 scripts/check_data_workflow.py
```

Keep `technical-design.md` , `evidence-matrix.csv` , the generator, draw.io source and the dated verification record together. When a requirement changes, update its acceptance criterion and evidence row before claiming it complete.

Appendix · Complete evidence matrix

Each row identifies a need, an observable acceptance criterion, implementation evidence, and the remaining limit. Paths are relative to the LTEF project root. Named tests require a dated execution record before being described as passing.

BR-001

Implemented · verified

Know the evidence each metric needs

Acceptance: Authenticated users can inspect fields, applicability and collection guidance for all 18 definitions.

Implementation

Metric catalog + guided choice and preparation screens

File / function

ltef/requirements.py:data_requirements; ltef/web/data.js:chooseDataMetrics, prepareData

Test / evidence

tests/test_requirements.py:test_catalog_covers_executable_metrics_and_supplied_evidence; scripts/check_data_workflow.py

Gap / next action

Catalog must evolve with every scorer change.

Responsible role

Product / study lead

BR-002

Implemented · verified

Start with compatible files

Acceptance: Researcher/Admin downloads a valid ZIP filtered to selected sector and metric IDs.

Implementation

Synthetic three-file starter + README and correct prompt fingerprints

File / function

ltef/requirements.py:starter_bundle; ltef/web/data.js:prepareData

Test / evidence

tests/test_requirements.py:test_templates_are_sector_scoped_valid_and_explicitly_synthetic;
test_template_selection_is_exact_and_invalid_choices_rejected; scripts/check_data_workflow.py

Gap / next action

Starter is synthetic; it is not a representative dataset or collection adapter.

Responsible role

Dataset author

BR-003

Implemented · verified

Supply a consistent data contract

Acceptance: Reject duplicate observations/JSON keys, unknown cases, invalid numbers and malformed scope/policy fields before saving.

Implementation

JSON/JSONL parsing plus canonical schema validation

File / function

ltef/schema.py; ltef/requirements.py:load_upload; ltef/webapp.py:Workspace.run

Test / evidence

tests/test_ltef.py:test_duplicate_observation_rejected, test_unknown_case_rejected,
test_nonfinite_measurement_and_boolean_token_count_rejected, test_duplicate_json_keys_rejected;
tests/test_webapp.py:test_invalid_inputs_do_not_create_runs

Gap / next action

No form editor or CSV conversion; users provide schema-compliant files.

Responsible role

Collection engineer

BR-004

Implemented · verified

Bind observations to the prompts tested

Acceptance: Every record references a known case/system/trial and exact canonical messages digest.

Implementation

Record linkage and SHA-256 match validation

File / function

ltef/schema.py:digest, validate_bundle

Test / evidence

tests/test_ltef.py:test_wrong_prompt_rejected; tests/test_webapp.py:test_prompt_hash_mismatch_rejected_and_atomic

Gap / next action

Hashes detect internal mismatch; provider capture remains unverified.

Responsible role

Collection engineer

BR-005

Implemented · verified

Check readiness without polluting history

Acceptance: Preflight returns valid/invalid structure and per-metric coverage without saving; replacing a file removes run readiness.

Implementation

In-memory readiness + guarded browser payload; run revalidates

File / function

ltef/requirements.py:readiness; ltef/webapp.py:Handler.do_POST; ltef/web/data.js:validatePreparedData, runPreparedData

Test / evidence

tests/test_requirements.py:test_readiness_does_not_persist_and_actual_run_revalidates,
test_readiness_distinguishes_missing_failed_and_inapplicable; scripts/check_data_workflow.py

Gap / next action

Structural validity does not validate sampling, annotations or outcomes.

Responsible role

Researcher

BR-006

Implemented · verified

Keep metrics scoped to a study

Acceptance: Profile fixes sector/metrics/grouping/minimum cases and rejects calibrated thresholds.

Implementation

Versioned profiles; selected sector case set

File / function

ltef/schema.py:validate_profile; ltef/engine.py:evaluate

Test / evidence

tests/test_ltef.py:test_no_calibrated_thresholds_implied;
tests/test_requirements.py:test_template_selection_is_exact_and_invalid_choices_rejected

Gap / next action

Sector naming alone does not establish domain validity.

Responsible role

Study lead

BR-007

Implemented · verified

Measure deterministic answer accuracy and confidence

Acceptance: Exact/numeric matching rejects extraneous numeric answers and nonfinite values; calibration uses reference correctness.

Implementation

Reference matching and Brier scorer

File / function

ltef/metrics.py:accuracy, score

Test / evidence

tests/test_ltef.py:test_numeric_reference_does_not_accept_extraneous_answer_or_nan,
test_case_normalization_supports_unicode_and_whitespace, test_demo_has_known_outcomes_and_missingness

Gap / next action

No semantic answer judge; confidence validity must be established externally.

Responsible role

Evaluation engineer

BR-008

Implemented · verified

Avoid treating missing review as a clean result

Acceptance: Annotation identity/method/rubric is required; absent or empty claim reviews remain unmeasured.

Implementation

Annotation validation and supported-claim calculation

File / function

ltef/schema.py:validate_bundle; ltef/metrics.py:score

Test / evidence

tests/test_ltef.py:test_annotation_requires_provenance, test_empty_claims_never_pass, test_missing_annotations_never_pass

Gap / next action

Rubric quality, claim completeness and reviewer agreement are outside scoring.

Responsible role

Domain adjudicator

BR-009

Implemented · verified

Use clinical metrics only with clinical evidence

Acceptance: Healthcare cases require explicitly supplied clinical review; no review yields zero measured observations, not zero clinical errors.

Implementation

Healthcare eligibility and major-error indicator

File / function

ltef/metrics.py:score; ltef/requirements.py:readiness

Test / evidence

tests/test_requirements.py:test_missing_reviews_do_not_become_zero_error_results; scripts/check_data_workflow.py

Gap / next action

Clinical validity and clinician agreement have not been established.

Responsible role

Clinical study lead

BR-010

Implemented · verified

Measure defined safety/privacy behavior

Acceptance: Apply canary, attack, benign and escalation metrics only to scoped cases with needed responses/labels.

Implementation

Narrow privacy probe and supplied safety outcome scorers

File / function

ltef/metrics.py:score; ltef/schema.py:validate_cases

Test / evidence

tests/test_ltef.py:test_demo_has_known_outcomes_and_missingness, test_missing_annotations_never_pass, test_trace_privacy_scope_excludes_cases_without_canoaries

Gap / next action

These probes do not establish comprehensive security, privacy or safe use.

Responsible role

Safety evaluator

Account for resources consumed by failures

Acceptance: Supplied nonnegative latency/cost/token values also contribute on failed observations.

Implementation

Resource scorers independent of successful response scoring

File / function

ltef/metrics.py:score

Test / evidence

tests/test_ltef.py:test_failed_calls_still_count_measured_resources,
test_nonfinite_measurement_and_boolean_token_count_rejected

Gap / next action

No live billing, tokenizer or wall-clock capture in the web app.

Responsible role

Collection engineer

Inspect multi-agent behavior with explicit evidence

Acceptance: Separate unauthorized attempts from executions; require complete traces for trace metrics; validate receipt ordering.

Implementation

Seven agent metrics + strict trace event contract

File / function

ltef/metrics.py:score; ltef/schema.py:validate_trace

Test / evidence

tests/test_ltef.py:test_unauthorized_attempt_distinct_from_execution, test_incomplete_trace_is_missing_not_safe,
test_no_handoffs_not_one_hundred_percent, test_unreceived_handoff_scores_zero, test_wrong_handoff_recipient_rejected,
test_future_parent_and_duplicate_receive_rejected, test_propagation_needs_source_trace

Gap / next action

Task success and propagation labels supplied externally; completeness is asserted, not independently proven.

Responsible role

Orchestration engineer

Expose missingness and case-weighted uncertainty

Acceptance: Retain all declared trial slots; report state counts/coverage; average trials within cases then equally weight cases.

Implementation

Case-weighted summaries and deterministic case bootstrap

File / function

ltef/statistics.py:summarize; ltef/engine.py:evaluate

Test / evidence

tests/test_ltef.py:test_case_weighting_does_not_overweight_repeated_trials, test_all_failed_system_remains_in_report, test_demo_has_known_outcomes_and_missingness

Gap / next action

Minimum 20 is a fixture reporting setting, not a power analysis; study dependence needs consideration.

Responsible role

Statistician / study lead

Compare equivalent observations

Acceptance: Pair baseline/candidate by case and trial within the same report; report unpaired observations and direction.

Implementation

Paired comparison and bootstrap of case-mean deltas

File / function

ltef/engine.py:compare; ltef/web/app.js:comparePage

Test / evidence

tests/test_ltef.py:test_comparison_pairs_trial_ids_before_averaging, test_comparison_direction_and_pairs; scripts/check_webapp.py

Gap / next action

No causal superiority claim, multiple-testing correction or cross-report comparison.

Responsible role

Study lead

Inspect variation across populations

Acceptance: Group and pair outputs remain interpretable; missing/small group evidence does not imply fairness.

Implementation

Named subgroup summaries and exact controlled-pair agreement

File / function

ltef/engine.py:evaluate; ltef/schema.py:validate_cases

Test / evidence

tests/test_ltef.py:test_unlabeled_groups_and_small_groups_not_fairness_pass, test_invalid_counterfactual_pair_rejected

Gap / next action

Exact output agreement is not demographic fairness or semantic invariance.

Responsible role

Dataset / fairness reviewer

Retain reproducible derived results

Acceptance: Saved/exported reports contain definitions, profile, state counts, limitations and input/code fingerprints.

Implementation

Structured JSON report, Markdown export, SQLite persistence

File / function

ltef/engine.py:evaluate, code_digest; ltef/reporting.py:render; ltef/webapp.py:Workspace.save

Test / evidence

tests/test_webapp.py:test_seeds_three_sectors_and_persists_without_reseeding,
test_upload_runs_actual_engine_and_omits_raw_content; scripts/check_webapp.py

Gap / next action

Source inputs and full runtime environment must be retained separately.

Responsible role

Researcher

Sign in with local accounts

Acceptance: First concurrent registration creates exactly one Admin; later accounts default Viewer; password policy is enforced.

Implementation

Salted PBKDF2 account store and atomic registration

File / function

ltef/auth.py:Accounts.register, Accounts.login

Test / evidence

tests/test_auth.py:test_concurrent_first_registration_creates_exactly_one_admin,
test_password_policy_and_duplicate_account_do_not_change_users,
test_wrong_password_and_unknown_account_use_same_response

Gap / next action

No email verification, recovery, SSO or restricted invitations.

Responsible role

Workspace administrator

Enforce role access on the server

Acceptance: Viewer receives six metric values; run/compare/export require Researcher; role changes apply to active sessions.

Implementation

Permission guards and deep report filtering

File / function

ltef/auth.py:permissions, filter_report; ltef/webapp.py:Handler.require_user

Test / evidence

tests/test_auth.py:test_viewer_values_are_filtered_and_restricted_actions_fail,
test_role_changes_apply_to_existing_sessions_and_preserve_last_admin;
tests/test_requirements.py:test_permissions_guard_readiness_and_template_downloads; scripts/check_webapp.py

Gap / next action

Product roles are not credentials for domain adjudication.

Responsible role

Workspace administrator

BR-019

Implemented · verified

Keep evaluations private to their creator

Acceptance: New reports are private even from other Admins; only original seed examples are shared.

Implementation

Ownership predicates and explicit shared_sample flag

File / function

ltef/webapp.py:Workspace.get, Workspace.listing, Workspace.save

Test / evidence

tests/test_auth.py:test_private_runs_are_inaccessible_even_to_other_admin,
test_only_original_seeds_shared_and_legacy_runs_assigned_first_admin

Gap / next action

Operating-system access can bypass application-level ownership.

Responsible role

Workspace administrator

BR-020

Implemented · verified

Bound and revoke local sessions

Acceptance: Login rotates session, logout revokes, expiry survives restart; host/origin/CSRF checks reject invalid requests.

Implementation

Cookie session, token hashes, current-role lookup, request checks

File / function

ltef/auth.py:Accounts; ltef/webapp.py:Handler.trusted, Handler.do_POST

Test / evidence

tests/test_auth.py:test_initial_owner_login_rotates_session_and_logout_revokes,
test_passwords_saltd_and_sessions_survive_restart_until_expiry, test_csrf_origin_and_host_enforced,
test_login_attempts_are_bounded

Gap / next action

HTTP loopback deployment only; no production identity/security certification.

Responsible role

Security reviewer

Bound expensive requests

Acceptance: Reject excessive input bytes/case expansion and serialize evaluations/preflight with a shared lock.

Implementation

Request cap and upload limits before scoring

File / function

ltef/requirements.py:load_upload; ltef/webapp.py:Handler.do_POST

Test / evidence

tests/test_webapp.py:test_expansion_limit_applies_before_evaluation;
tests/test_requirements.py:test_errors_limits_csrf_and_download_contract

Gap / next action

No job queue, benchmarked throughput target or cancellation API.

Responsible role

Application engineer

Minimize saved sensitive content

Acceptance: Exclude raw responses, canary text, trace bodies and reviewer evidence from derived reports.

Implementation

Derived-report construction; raw inputs remain outside SQLite

File / function

ltef/engine.py:evaluate; ltef/webapp.py:Workspace.save

Test / evidence

tests/test_ltef.py:test_reports_do_not_include_raw_secret_or_output;
tests/test_webapp.py:test_upload_runs_actual_engine_and_omits_raw_content

Gap / next action

Metadata may still contain sensitive strings; no automatic redaction/retention/deletion workflow.

Responsible role

Data steward

BR-023

Implemented · verified

Navigate and provide evidence simply

Acceptance: Desktop and 390px mobile flows support auth, metric discovery, three intake steps, errors and results without page overflow.

Implementation

Responsive native frontend with keyboard/dialog handling

File / function

ltef/web/app.js; ltef/web/data.js; ltef/web/style.css

Test / evidence

scripts/check_webapp.py; scripts/check_data_workflow.py; runs/web/data-choose,requirements,mobile,check,check-mobile.png

Gap / next action

Formal accessibility audit and broader device/user testing remain open.

Responsible role

Product designer

BR-024

Implemented · documentation reviewed

Trace BRD requirements to evidence

Acceptance: Every requirement identifies an observable criterion, code, verification status and limitation.

Implementation

28-row matrix, verification record and technical design

File / function

docs/design/evidence-matrix.csv; docs/design/verification-record.md; docs/design/technical-design.md

Test / evidence

Named-test existence check and generated HTML matrix row count; source review

Gap / next action

Maintain on every material change; software evidence does not demonstrate societal impact.

Responsible role

Technical lead

BR-025

Implemented · artifact verified

Review and edit architecture diagrams

Acceptance: Provide five native draw.io pages with shapes/connectors and matching SVG previews.

Implementation

Architecture, data readiness, sequence, logical model and evidence traceability

File / function

scripts/build_design_diagrams.py; docs/design/LTEF-technical-design.drawio

Test / evidence

XML structure/reference validation; SVG/HTML visual review recorded in verification-record.md

Gap / next action

MCP configuration

Responsible role

initialization and tool discovery verified. Restart Codex to load the connector; diagram tool execution is not claimed.

BR-026

Partial

Capture live model observations

Acceptance: A provider adapter records responses/errors and resources without leaking reference labels into the request.

Implementation

Callback collection boundary exists; production adapter not implemented

File / function

ltef/collection.py:collect

Test / evidence

tests/test_ltef.py:test_collector_hides_reference_and_retains_sanitized_failures, test_no_synthetic_labels_from_adapter

Gap / next action

Implement provider-specific capture, billing basis, model version capture and retry protocol.

Responsible role

Collection engineer

BR-027

Planned

Substantiate real sector and multi-agent performance

Acceptance: Frozen representative cases, validated rubric, reviewer agreement, justified sampling and reproducible real captures support claimed findings.

Implementation

Research protocol and proposed 36-month plan

File / function

docs/measurement-protocol.md; docs/research-traceability.md; docs/roadmap-36-months.md

Test / evidence

No completed real-provider multi-sector validation study is claimed

Gap / next action

Collect independent evidence; validate new agent metrics and comparisons.

Responsible role

Research lead

BR-028

Planned

Operate as a production hosted service

Acceptance: Authenticated public deployment meets agreed identity, isolation, recovery, observability and availability criteria.

Implementation

Local deployment boundary explicitly documented

File / function

docs/design/technical-design.md: architecture, security and backlog

Test / evidence

No production certification, load test or recovery evidence is claimed

Gap / next action

Define production SLOs, HTTPS, identity/recovery, isolation, audit, backups, retention and operations.

Responsible role

Platform lead

Prepared from the local LTEF source and research protocol. Keep source inputs and dated verification logs alongside this review package. Editable diagrams: LTEF-technical-design.drawio.